



Book Cover Judger

ML Guild Apprentice Capstone – Grant Tesdahl

September 24, 2020

Contents

- 1 About Me**
- 2 Problem Statement**
- 3 Methodology**
- 4 Lessons Learned & Next Steps**



About Me



Grant Tesdahl

Analyst – CBO CCG
Minneapolis

Education

B.S. in Industrial Engineering from the University of Wisconsin-Madison



Project Experience

Data Modernization at Wells Fargo (rolling off next Friday!)



Interests

Reading, skiing, bouldering, soccer, & Freakonomics podcast



ML Journey

**Before
Deloitte:**

Data Rich Internships at UW Health & ExxonMobil...
...but without the tools to uncover the best insights



Data Mining & Machine Learning Course...
...where I realized that ML enables those insights

**At
Deloitte:**

Building **business cases** for predictive models
Banking Profitability Insights & Diabetes Mobile App



ML Guild Apprentice Program
Training + Capstone



Building **predictive models**
Cognitive Guided Tour: Revenue Cycle Modeling

Problem Statement


People Judge Books By Their Covers.

In fact, *Publishers Weekly* found that 75% of booksellers feel that the **cover's design is the most important element** in book promotion

How can publishers ensure a given cover will maximize a book's popularity?

Methodology



Using book rating data and book cover images from Goodreads, I built a **CNN that predicts book popularity**




1. Pull book review data & book cover images
2. **Calculate potential target values** related to a book's average rating and rating count

1. Train a **CNN for each potential target value**
2. Compare models and pick the best target variable
3. Run additional training

1. Design and code an interactive website
2. **Integrate the model** into the website



Data: Acquisition



Two types of data required for this computer vision task:

Book Cover Images + Book Rating Data

Approach

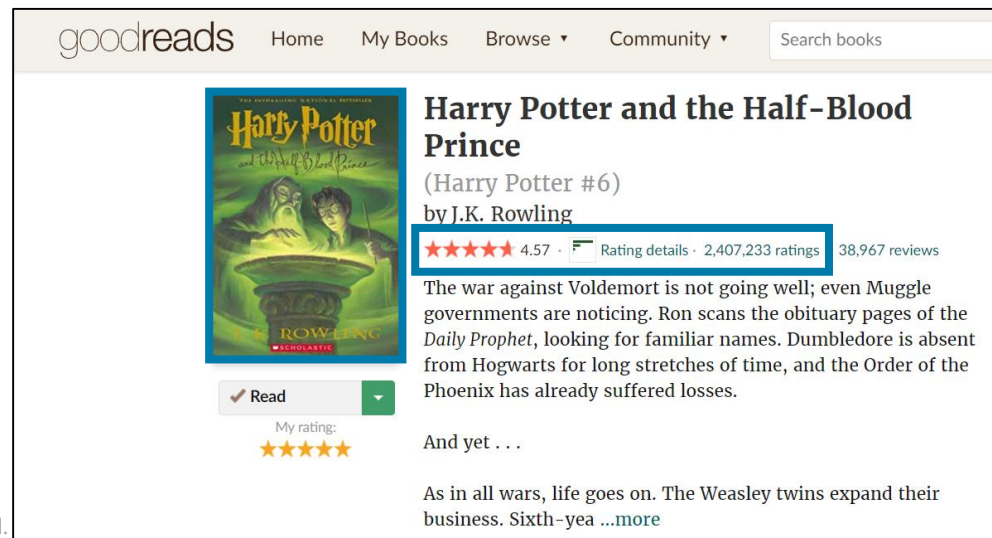
Captured 10k of each via Goodreads, which assigns a **bookID** to each book

Book Rating Data: **Average Rating** (indicates quality) & **Rating Count** (indicates popularity)

bookID = 1




<https://www.goodreads.com/book/show/1>



Example


Data: Transformation



Book Cover Images


Transformed images due to unequal aspect ratios

Image Aspect Ratio




Set aspect ratio to 1.5; Padded images with a reflection; Normalized pixel values

Original Image (Ratio = 1)




Transformed Image (Ratio = 1.5)




Book Rating Data

Transformed data to create 6 potential targets (4 regression, 2 classification)


Average Rating




Min-Max Norm of Rating



Rating Count



Log of Rating Count



Classification Segment Ranges

Segment		Segment	
Average Rating		Rating Count	
High	4.07 to 5.00	High	2,627 to 4,597,666
Medium	3.84 to 4.07	Medium	223 to 2,627
Low	0.00 to 3.84	Low	0 to 223

Modeling: Initial Training

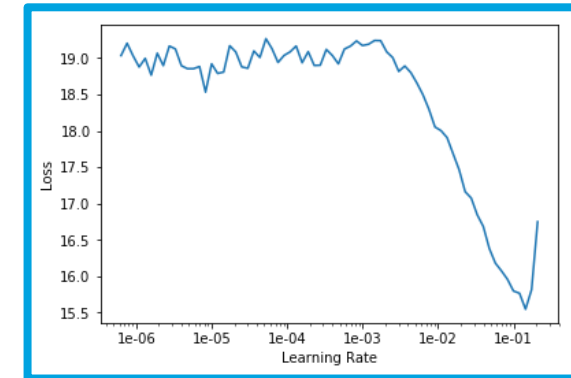


For Each Target Variable:

- Randomly split 20% of the samples for the **validation set**
- Create **mini-batches with 12 images** of size 192x128 (determined via experimentation)
- Create **CNN based on the ResNet34 Architecture** to start with a pre-trained model rather than from scratch
- Identify a **learning rate** to ensure rapid convergence
- Train for **5-8 epochs** depending on whether the loss continued decreasing
- Save validation set **predictions** to compare target variables



```
learn = cnn_learner(data, models.resnet34, metrics=root_mean_squared_error)
```



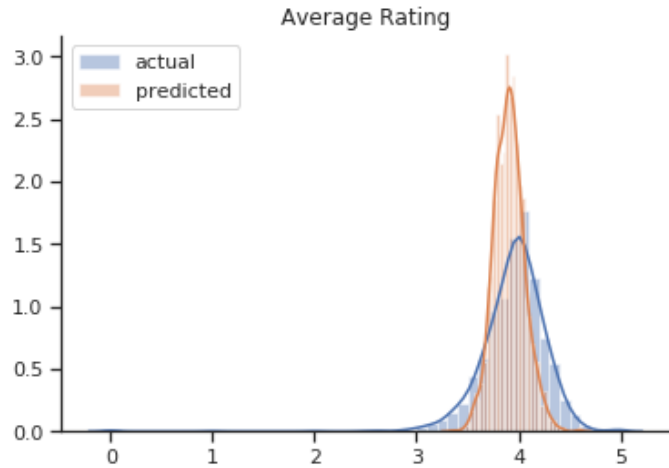
	average_rating_actual	average_rating_predicted	average_rating_error
0	3.19	3.869111	-0.679111
1	3.93	3.913237	0.016763
2	3.73	3.861418	-0.131418
3	3.68	3.951411	-0.271411
4	3.78	3.879087	-0.099087

Modeling: Target Variable Selection

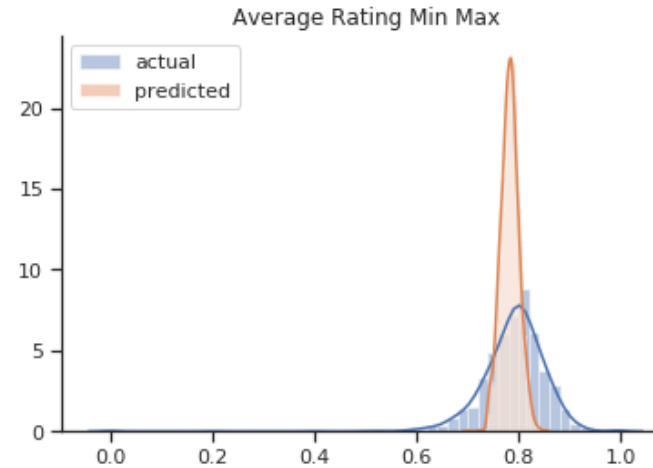


Rating Count Classification looks the least bad (42% accuracy); regression models center heavily around the mean

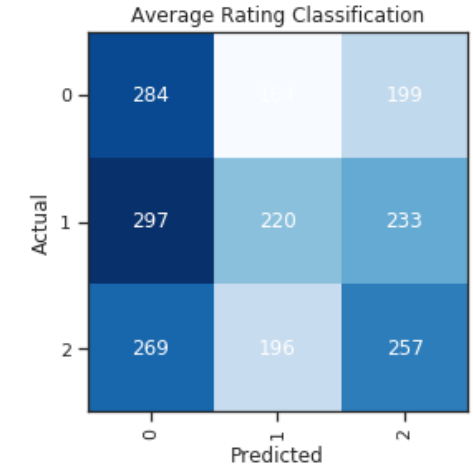
Regression: Raw Values



Regression: Transformed Values




Classification: High, Medium, & Low




Average Rating

Rating Count


Ratings Count



Ratings Count Log



Ratings Count Classification



Modeling: Additional Training



Enhancing the Classification Model:

Unfreeze ResNet34 Layers &
Train for 20 Epochs




Recreate Dataset with
Larger Images



Refreeze & Train Last Layers
for 3 Epochs


Final Architecture & Results:



44% Accuracy on the Validation Set
(Compared to 33% Baseline; $P\text{-Value} < 0.0001$)

	0	1	2
Actual 0	338	183	176
Actual 1	230	248	264
Actual 2	128	198	354
	0	1	2
Predicted			

Deployment



www.book-cover-judger.com

Lessons Learned & Next Steps

LESSONS LEARNED

1

Programming in Python

- My previous programming experience mainly consisted of DataCamp exercises, so a full project in Jupyter was a constant learning experience

2

Deep Learning

- The Neural Net lecture during the Bootcamp flew way over my head, so I really enjoyed having extra time to learn about CNNs, transfer learning, and other details

3

Importance of Experimentation

- Looking back, I used the recommended hyperparameters and architecture structures too frequently. In future projects, I plan to experiment with different combinations so I can see what works best for my specific use case

NEXT STEPS

1

Determine Most Important Features

- Extract and visualize the feature maps of each CNN block
- Translate activations into importance

2

Generate Book Covers with a GAN

- Collect more book covers & subset high popularity ones
- Train GAN on subset

3

Build in PyTorch from Scratch

- Leverage complex architectures (e.g. ResNet101) & customize
- Transition to an ordinal multi-label problem

Thank you!

